

plot_hog

October 26, 2020

```
[ ]: %matplotlib inline
```

1 Histogram of Oriented Gradients

The Histogram of Oriented Gradient (HOG) feature descriptor is popular for object detection [1].

In the following example, we compute the HOG descriptor https://en.wikipedia.org/wiki/Histogram_of_oriented_gradients and display a visualisation.

1.1 Algorithm overview

Compute a Histogram of Oriented Gradients (HOG) by

1. (optional) global image normalisation
2. computing the gradient image in x and y
3. computing gradient histograms
4. normalising across blocks
5. flattening into a feature vector

The first stage applies an optional global image normalisation equalisation that is designed to reduce the influence of illumination effects. In practice we use gamma (power law) compression, either computing the square root or the log of each color channel. Image texture strength is typically proportional to the local surface illumination so this compression helps to reduce the effects of local shadowing and illumination variations.

The second stage computes first order image gradients. These capture contour, silhouette and some texture information, while providing further resistance to illumination variations. The locally dominant color channel is used, which provides color invariance to a large extent. Variant methods may also include second order image derivatives, which act as primitive bar detectors - a useful feature for capturing, e.g. bar like structures in bicycles and limbs in humans.

The third stage aims to produce an encoding that is sensitive to local image content while remaining resistant to small changes in pose or appearance. The adopted method pools gradient orientation information locally in the same way as the SIFT [2] feature. The image window is divided into small spatial regions, called “cells”. For each cell we accumulate a local 1-D histogram of gradient or edge orientations over all the pixels in the cell. This combined cell-level 1-D histogram forms the basic “orientation histogram” representation. Each orientation histogram divides the gradient angle range into a fixed number of predetermined bins. The gradient magnitudes of the pixels in the cell are used to vote into the orientation histogram.

The fourth stage computes normalisation, which takes local groups of cells and contrast normalises their overall responses before passing to next stage. Normalisation introduces better invariance to illumination, shadowing, and edge contrast. It is performed by accumulating a measure of local histogram “energy” over local groups of cells that we call “blocks”. The result is used to normalise each cell in the block. Typically each individual cell is shared between several blocks, but its normalisations are block dependent and thus different. The cell thus appears several times in the final output vector with different normalisations. This may seem redundant but it improves the performance. We refer to the normalised block descriptors as Histogram of Oriented Gradient (HOG) descriptors.

The final step collects the HOG descriptors from all blocks of a dense overlapping grid of blocks covering the detection window into a combined feature vector for use in the window classifier.

1.2 References

- .. [1] Dalal, N. and Triggs, B., “Histograms of Oriented Gradients for Human Detection,” IEEE Computer Society Conference on Computer Vision and Pattern Recognition, 2005, San Diego, CA, USA.
- .. [2] David G. Lowe, “Distinctive image features from scale-invariant keypoints,” International Journal of Computer Vision, 60, 2 (2004), pp. 91-110.

```
[ ]: #import all related packages
import matplotlib.pyplot as plt
import numpy as np
import cv2

from skimage.feature import hog
from skimage import data, exposure, transform, color, feature

[ ]: image = data.astronaut()

fd, hog_image = hog(image, orientations=8, pixels_per_cell=(16, 16),
                    cells_per_block=(1, 1), visualize=True, multichannel=True)

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(8, 4), sharex=True, sharey=True)

ax1.axis('off')
ax1.imshow(image, cmap=plt.cm.gray)
ax1.set_title('Input image')

# Rescale histogram for better display
hog_image_rescaled = exposure.rescale_intensity(hog_image, in_range=(0, 10))

ax2.axis('off')
ax2.imshow(hog_image_rescaled, cmap=plt.cm.gray)
ax2.set_title('Histogram of Oriented Gradients')
plt.show()
```

2 HOG in Action: A Simple Face Detector

Using these HOG features, we can build up a simple facial detection algorithm with any Scikit-Learn estimator. The steps are as follows:

1. Obtain a set of image thumbnails of faces to constitute “positive” training samples.
2. Obtain a set of image thumbnails of non-faces to constitute “negative” training samples.
3. Extract HOG features from these training samples.
4. Train a linear SVM classifier on these samples.
5. For an “unknown” image, pass a sliding window across the image, using the model to evaluate whether that window contains a face or not.
6. If detections overlap, combine them into a single window.

Let’s go through these steps and try it out:

reference: <https://jakevdp.github.io/PythonDataScienceHandbook/05.14-image-features.html>

2.0.1 1. Obtain a set of positive training samples

Let’s start by finding some positive training samples that show a variety of faces. We have one easy set of data to work with—the ‘Labeled Faces in the Wild dataset’, which can be downloaded by Scikit-Learn:

```
[ ]: from sklearn.datasets import fetch_lfw_people
faces = fetch_lfw_people()
positive_patches = faces.images
positive_patches.shape
```

2.0.2 2. Obtain a set of negative training samples

Next we need a set of similarly sized thumbnails which do not have a face in them. One way to do this is to take any corpus of input images, and extract thumbnails from them at a variety of scales. Here we can use some of the images shipped with Scikit-Image, along with Scikit-Learn’s PatchExtractor:

```
[ ]: imgs_to_use = ['camera', 'text', 'coins', 'moon',
                  'page', 'clock', 'immunohistochemistry',
                  'chelsea', 'coffee', 'hubble_deep_field']
images = [color.rgb2gray(getattr(data, name)())
          for name in imgs_to_use]
```

```
[ ]: from sklearn.feature_extraction.image import PatchExtractor

def extract_patches(img, N, scale=1.0, patch_size=positive_patches[0].shape):
    extracted_patch_size = tuple((scale * np.array(patch_size)).astype(int))
    extractor = PatchExtractor(patch_size=extracted_patch_size,
                              max_patches=N, random_state=0)
    patches = extractor.transform(img[np.newaxis])
    if scale != 1:
        patches = np.array([transform.resize(patch, patch_size)
```

```

        for patch in patches])

    return patches

negative_patches = np.vstack([extract_patches(im, 1000, scale)
                              for im in images for scale in [0.5, 1.0, 2.0]])
negative_patches.shape

```

We now have 30,000 suitable image patches which do not contain faces. Let's take a look at a few of them to get an idea of what they look like:

```

[ ]: fig, ax = plt.subplots(6, 10)
     for i, axi in enumerate(ax.flat):
         axi.imshow(negative_patches[500 * i], cmap='gray')
         axi.axis('off')

```

2.0.3 3. Combine sets and extract HOG features

Now that we have these positive samples and negative samples, we can combine them and compute HOG features. This step takes a little while, because the HOG features involve a nontrivial computation for each image:

```

[ ]: from itertools import chain

     X_train = np.array([feature.hog(im)
                        for im in chain(positive_patches,
                                        negative_patches)])
     y_train = np.zeros(X_train.shape[0])
     y_train[:positive_patches.shape[0]] = 1

     X_train.shape

```

2.0.4 4. Training a support vector machine

Next we use the tools we have been exploring in this chapter to create a classifier of thumbnail patches. For such a high-dimensional binary classification task, a Linear support vector machine is a good choice. We will use Scikit-Learn's LinearSVC, because in comparison to SVC it often has better scaling for large number of samples.

First, though, let's use a simple Gaussian naive Bayes to get a quick baseline:

```

[ ]: from sklearn.naive_bayes import GaussianNB
     from sklearn.model_selection import cross_val_score

     cross_val_score(GaussianNB(), X_train, y_train)

```

```

[ ]: from sklearn.svm import LinearSVC
     from sklearn.model_selection import GridSearchCV

```

```
grid = GridSearchCV(LinearSVC(), {'C': [1.0, 2.0]})
grid.fit(X_train, y_train)
grid.best_score_
```

```
[ ]: grid.best_params_
```

```
[ ]: model = grid.best_estimator_
model.fit(X_train, y_train)
```

2.0.5 5. Find faces in a new image

Now that we have this model in place, let's grab a new image and see how the model does. We will use one portion of the astronaut image for simplicity (see discussion of this in Caveats and Improvements), and run a sliding window over it and evaluate each patch:

For more images, see

<https://scikit-image.org/docs/dev/api/skimage.data.html?highlight=data#module-skimage.data>

```
[ ]: # Test images input for doing experiments

t1 = cv2.imread('images/1.jpg',0)
t1 = color.rgb2gray(t1)
t1 = transform.rescale(t1, 0.5)

t2 = cv2.imread('images/2.jpg',0)
t2 = color.rgb2gray(t2)
t2 = transform.rescale(t2, 0.5)

t3 = cv2.imread('images/3.jpg',0)
t3 = color.rgb2gray(t3)
t3 = transform.rescale(t3, 0.5)

t4 = cv2.imread('images/4.jpg',0)
t4 = color.rgb2gray(t4)
t4 = transform.rescale(t4, 0.5)

fig, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2)
ax1.axis('off')
ax1.imshow(t1, cmap=plt.cm.gray)
ax1.set_title('Test_image No. 1')

ax2.axis('off')
ax2.imshow(t2, cmap=plt.cm.gray)
ax2.set_title('Test_image No. 2')
```

```

ax3.axis('off')
ax3.imshow(t3, cmap=plt.cm.gray)
ax3.set_title('Test_image No. 3')

ax4.axis('off')
ax4.imshow(t4, cmap=plt.cm.gray)
ax4.set_title('Test_image No. 4')

plt.show()

```

```

[ ]: ### Changing the input of test images here ###
test_image = cv2.imread('images/1.jpg',0)

test_image = color.rgb2gray(test_image)
test_image = transform.rescale(test_image, 0.5)

def sliding_window(img, patch_size=positive_patches[0].shape,
                   istep=2, jstep=2, scale=1.0):
    Ni, Nj = (int(scale * s) for s in patch_size)
    for i in range(0, img.shape[0] - Ni, istep):
        for j in range(0, img.shape[1] - Nj, jstep):
            patch = img[i:i + Ni, j:j + Nj]
            if scale != 1:
                patch = transform.resize(patch, patch_size)
            yield (i, j), patch

indices, patches = zip(*sliding_window(test_image))
patches_hog = np.array([feature.hog(patch) for patch in patches])
patches_hog.shape

#Finally, we can take these HOG-featured patches and use our model
#to evaluate whether each patch contains a face:

labels = model.predict(patches_hog)
labels.sum()

## HOG features for display
fd, hog_image = feature.hog(test_image, orientations=8, pixels_per_cell=(16,16),
                             cells_per_block=(1, 1), visualize=True, multichannel=False)
hog_image_rescaled = exposure.rescale_intensity(hog_image, in_range=(0, 10))

fig, (ax1, ax2, ax3) = plt.subplots(1, 3, figsize=(16, 8), sharex=True,
                                     sharey=True)

```

```

ax1.axis('off')
ax1.imshow(test_image, cmap=plt.cm.gray)
ax1.set_title('Input image')

ax2.axis('off')
ax2.imshow(hog_image_rescaled, cmap=plt.cm.gray)
ax2.set_title('Histogram of Oriented Gradients')

ax3.axis('off')
ax3.imshow(test_image, cmap=plt.cm.gray)
ax3.set_title('Face detected')


Ni, Nj = positive_patches[0].shape
indices = np.array(indices)

for i, j in indices[labels == 1]:
    ax3.add_patch(plt.Rectangle((j, i), Nj, Ni, edgecolor='red',
                                alpha=0.3, lw=2, facecolor='none'))

plt.show()

```

All of the detected patches overlap and found the face in the image! Not bad for a few lines of Python.

Inline question 1: (ex, 1 - 4 images) De-
scribe the misclassification results that you see. Do they make sense?

[]: *### Answer here ###*